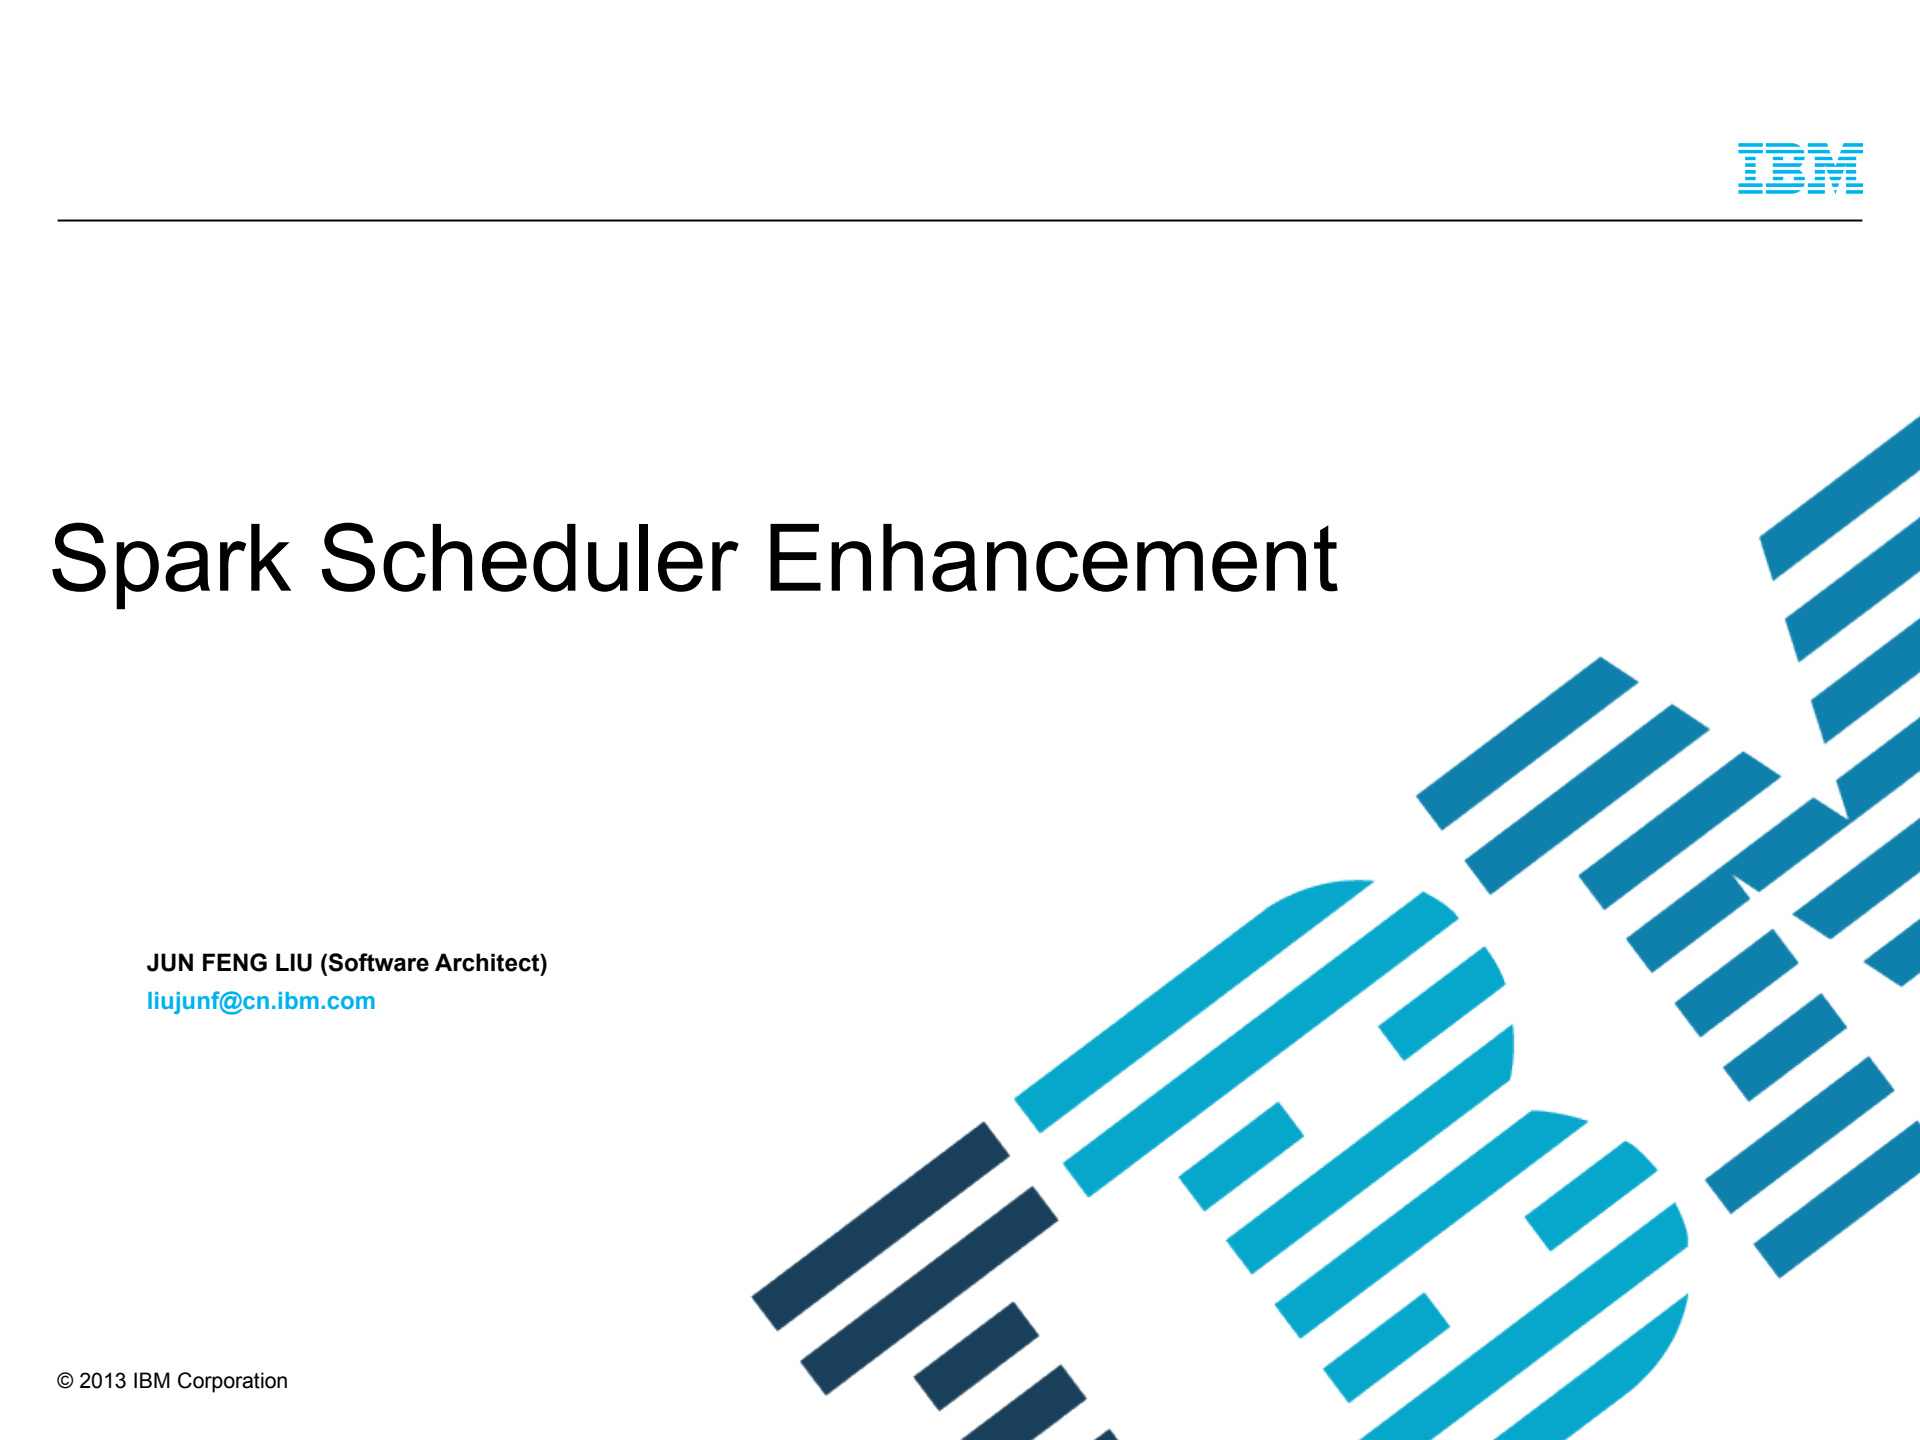

Spark Scheduler Enhancement

JUN FENG LIU (Software Architect)
liujunf@cn.ibm.com

Agenda

- **Background**
- **Spark on YARN**
- **Spark on Mesos**
- **Spark on IBM Platform Symphony**
- **Spark SQL or Stream**

Scheduling Process

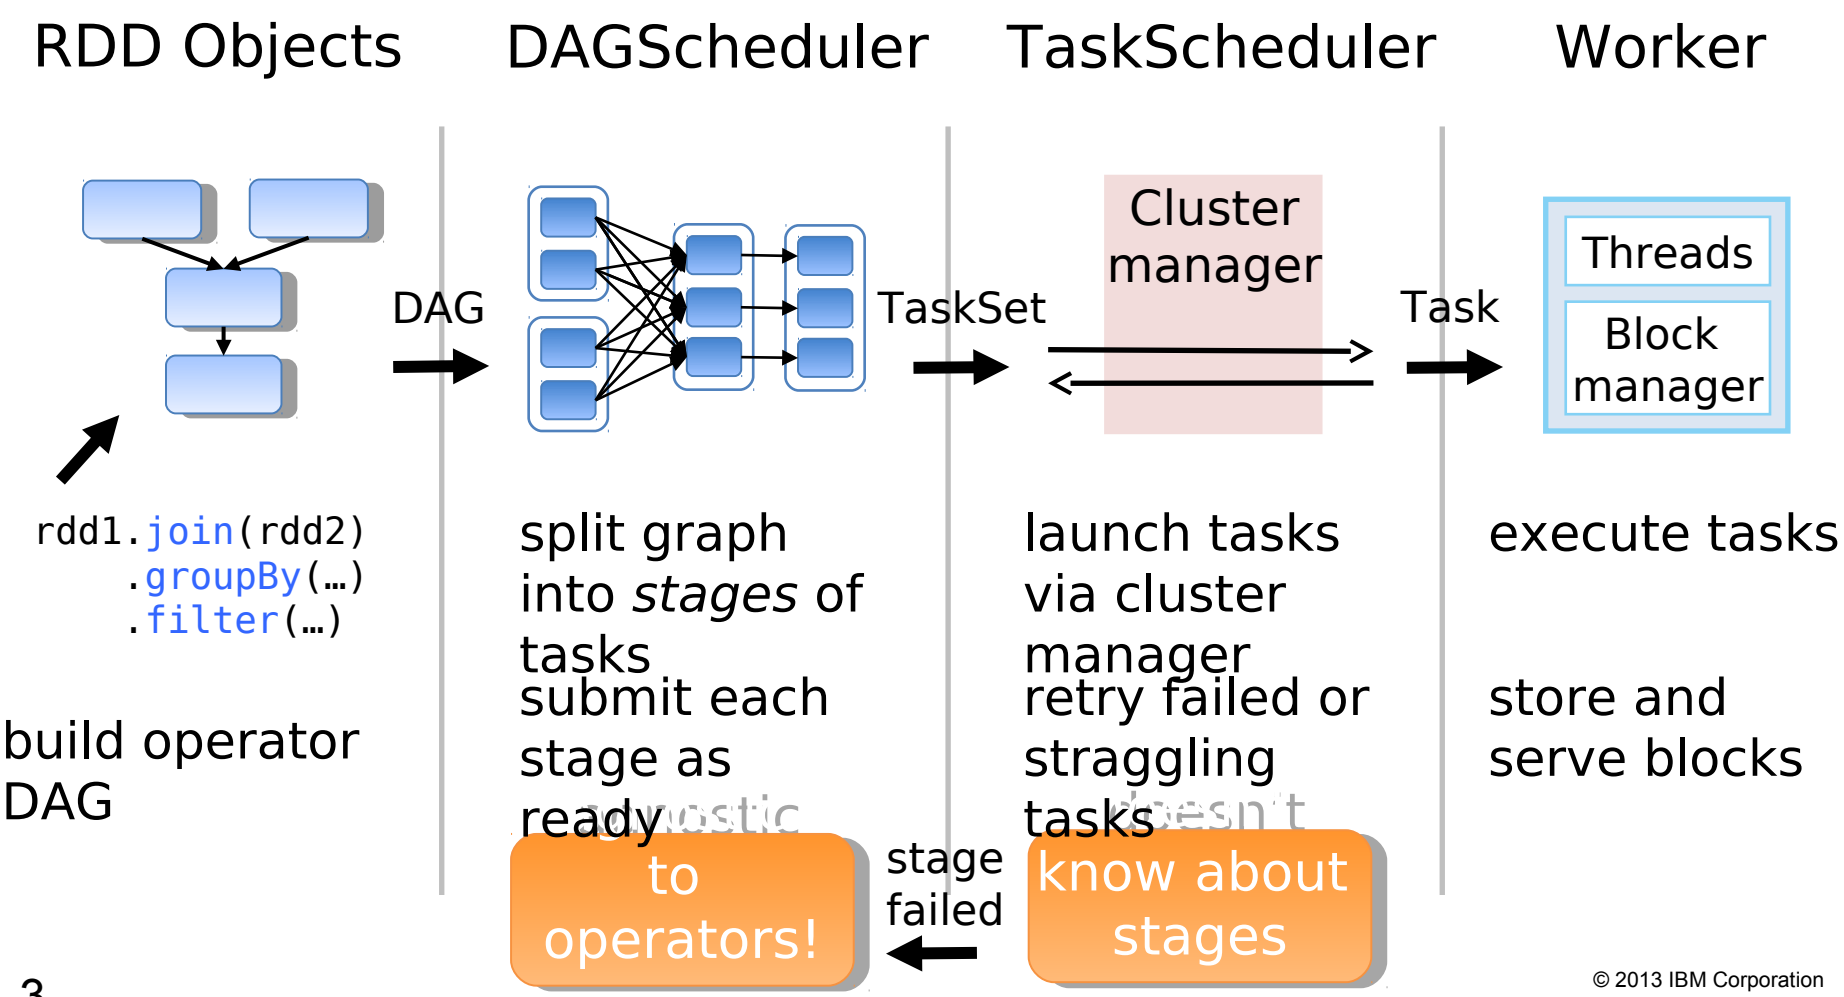

Event Flow

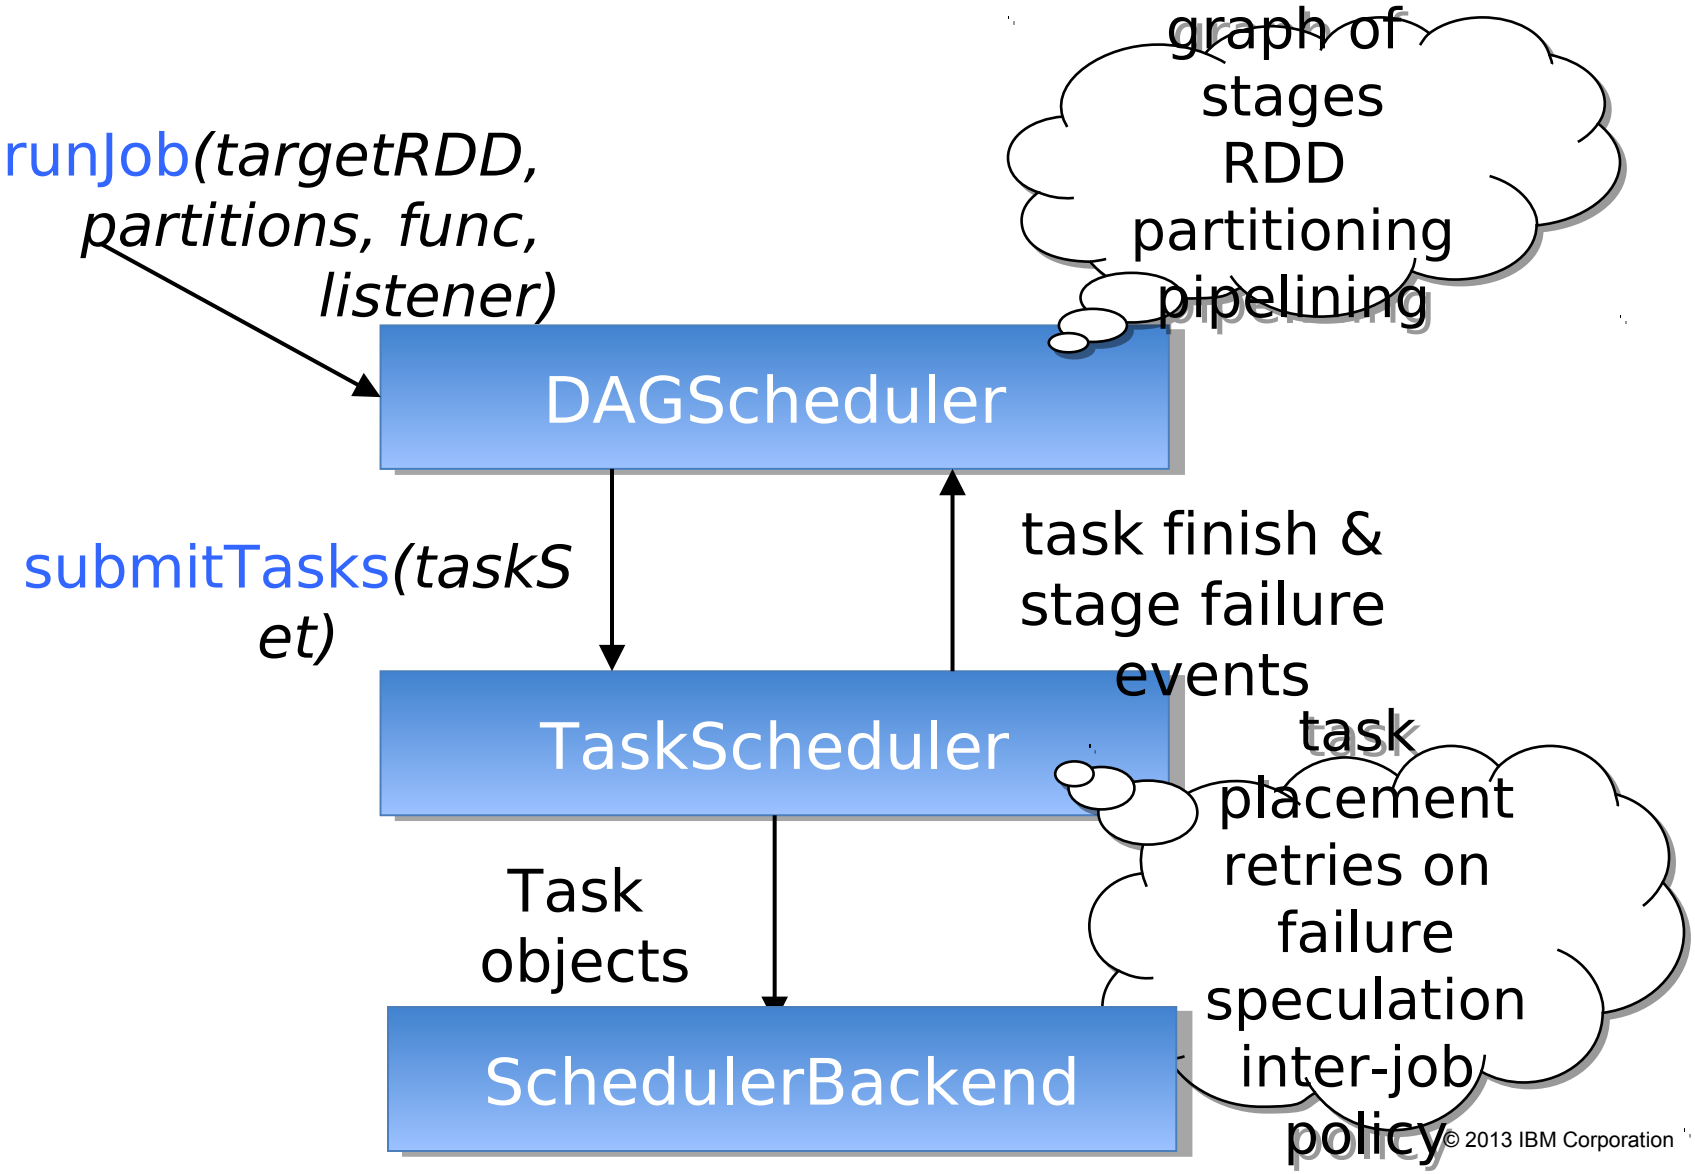

Agenda

- **Background**
- **Spark on YARN**
- **Spark on Mesos**
- **Spark on IBM Platform Symphony**
- **Spark SQL or Stream**

YARN

- **Coarse Model Scheduler**
  - Fix number Executor
  - Fix number Core
  - Fix number Memory

```
$ ./bin/spark-submit --class org.apache.spark.examples.SparkPi \
    --master yarn-cluster \
    --num-executors 3 \
    --driver-memory 4g \
    --executor-memory 2g \
    --executor-cores 1 \
    lib/spark-examples*.jar \
    10
```

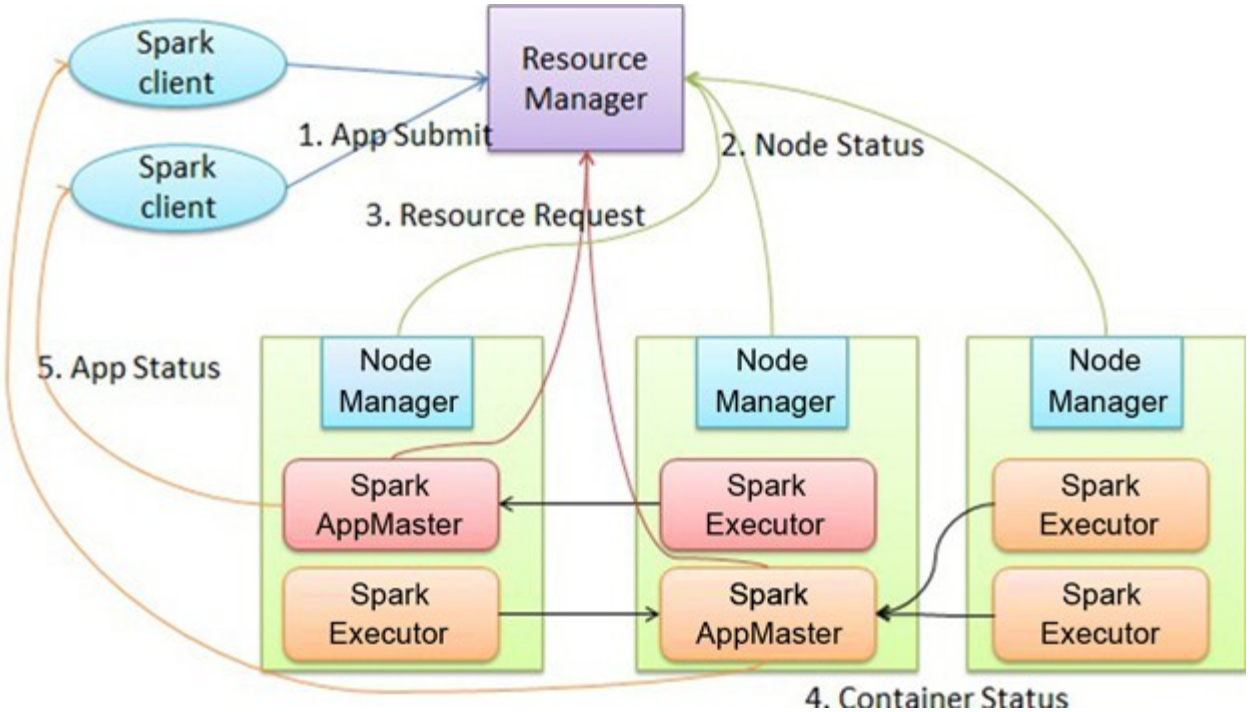

YARN

- **Data locality**
  - Rely on Driver Code (error prone)
  - Query the HDFS to get prefer hosts

```
val sparkConf = new SparkConf().setAppName("SparkHdfsLR")
    val inputPath = args(0)
    val conf = SparkHadoopUtil.get.newConfiguration()
    val sc = new SparkContext (sparkConf,
      InputFormatInfo.computePreferredLocations (
        Seq(new InputFormatInfo(conf,
  classOf[org.apache.hadoop.mapred.TextInputFormat],
  “hdfs://path/to/data.file”)) ))
```

Fine Grained on Mesos

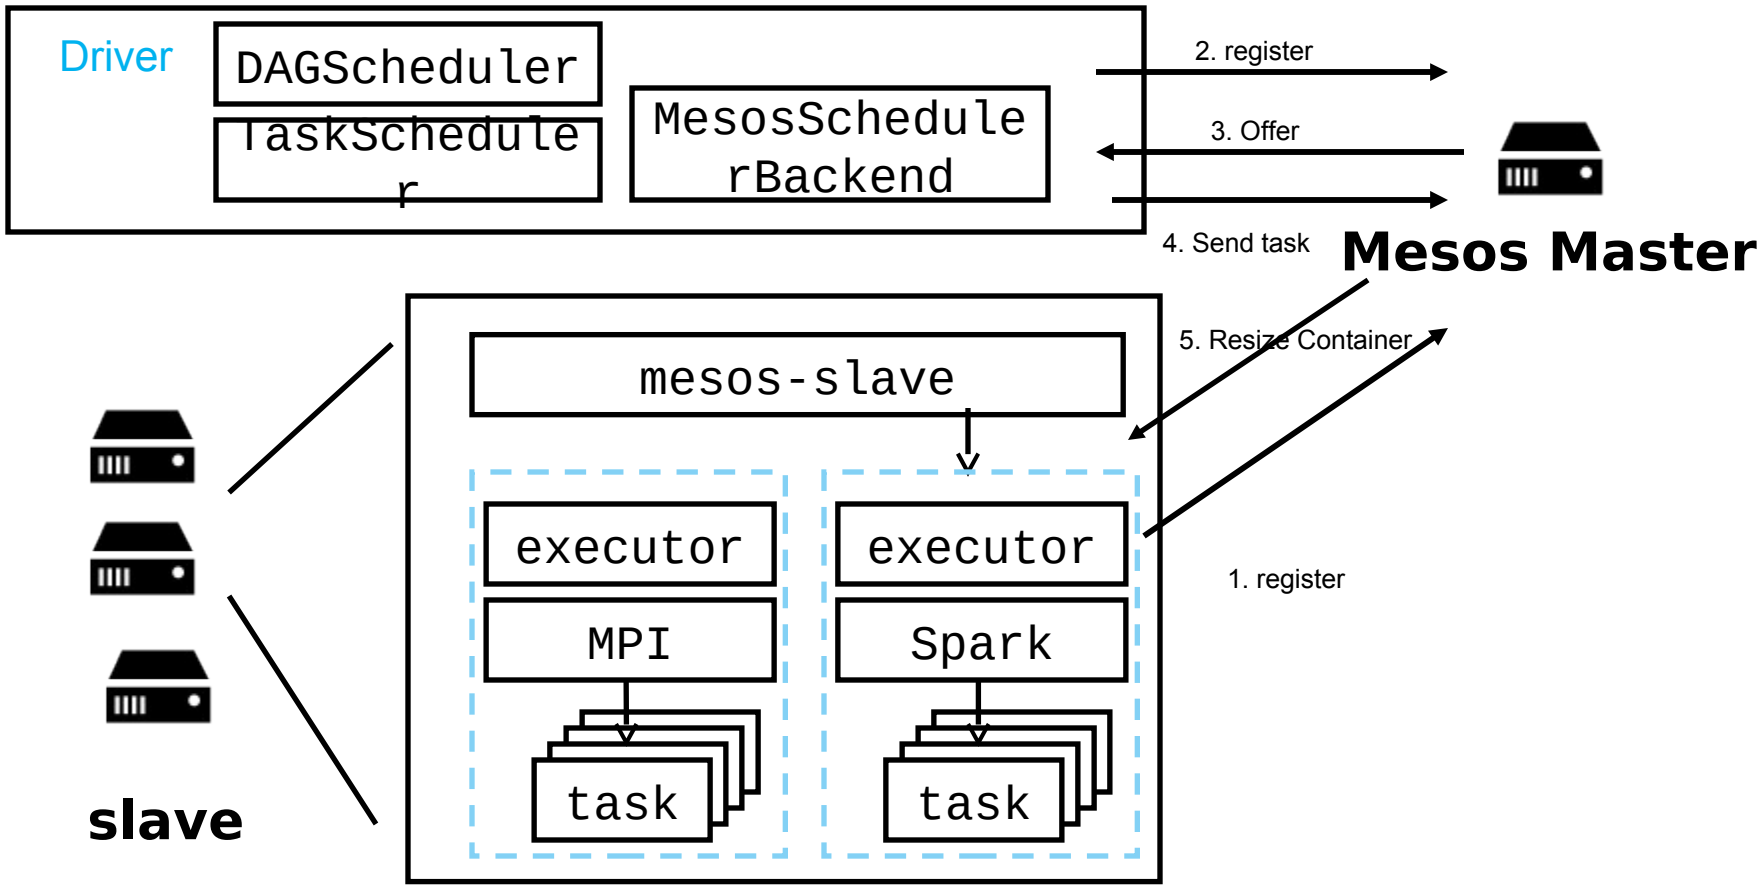

Fine Grained on Mesos

- **Resource Allocation**
  - Decided by Mesos DRF scheduler
  - Offer triggered when
    - Task Finished
    - Extra Slave add in
    - Application launch
  - Pessimistic offer
  - Revocable resources when task complete
  - Resource Share is good when task is short
  - No reclaim when task is long

Fine Grained on Mesos

- **Defer scheduling**
  - Spark can not choice preference host in Mesos model
  - Defer to schedule task if the data locality is not good with current offer
  - Defer scheduling in three level
    - spark.locality.wait.process
    - spark.locality.wait.node
    - spark.locality.wait.rack
-

Fine Grained on Mesos

pessimistic **optimistic**

offers made to different frameworks are disjoint

Mesos

all offers overlap with one another, thus causing frameworks to “compete” first-come-first-served

Google Omega

IBM Platform Symphony

Multiple users, applications and lines of business on a shared, heterogeneous, multi-tenant grid

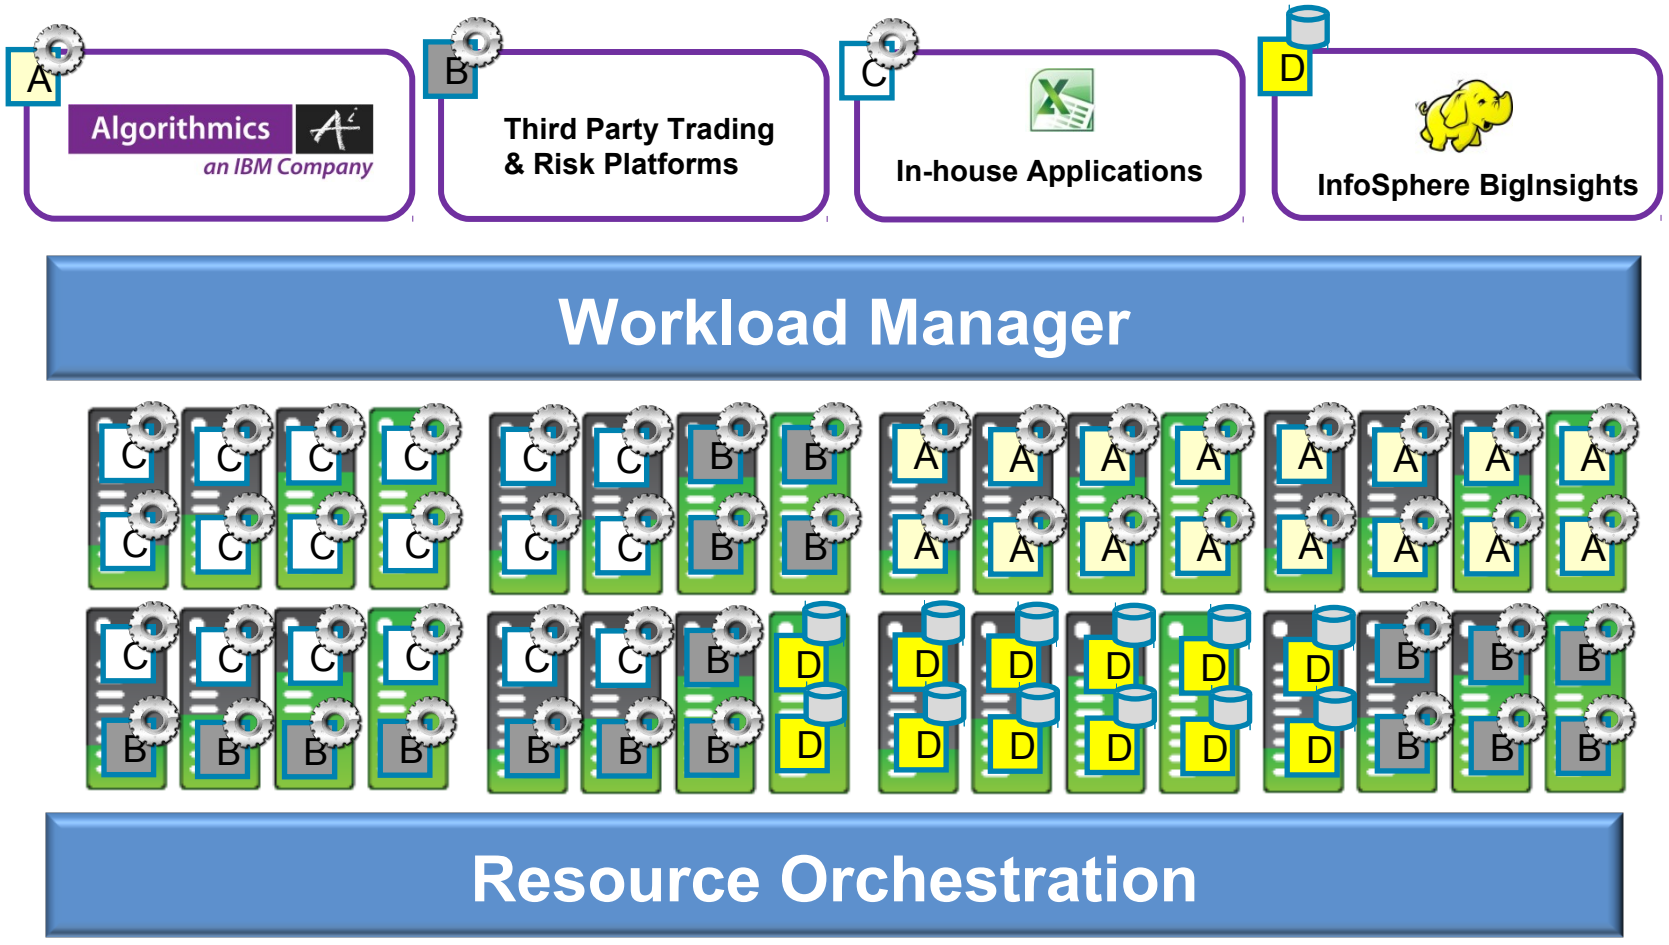

Application & Data Integration Architecture

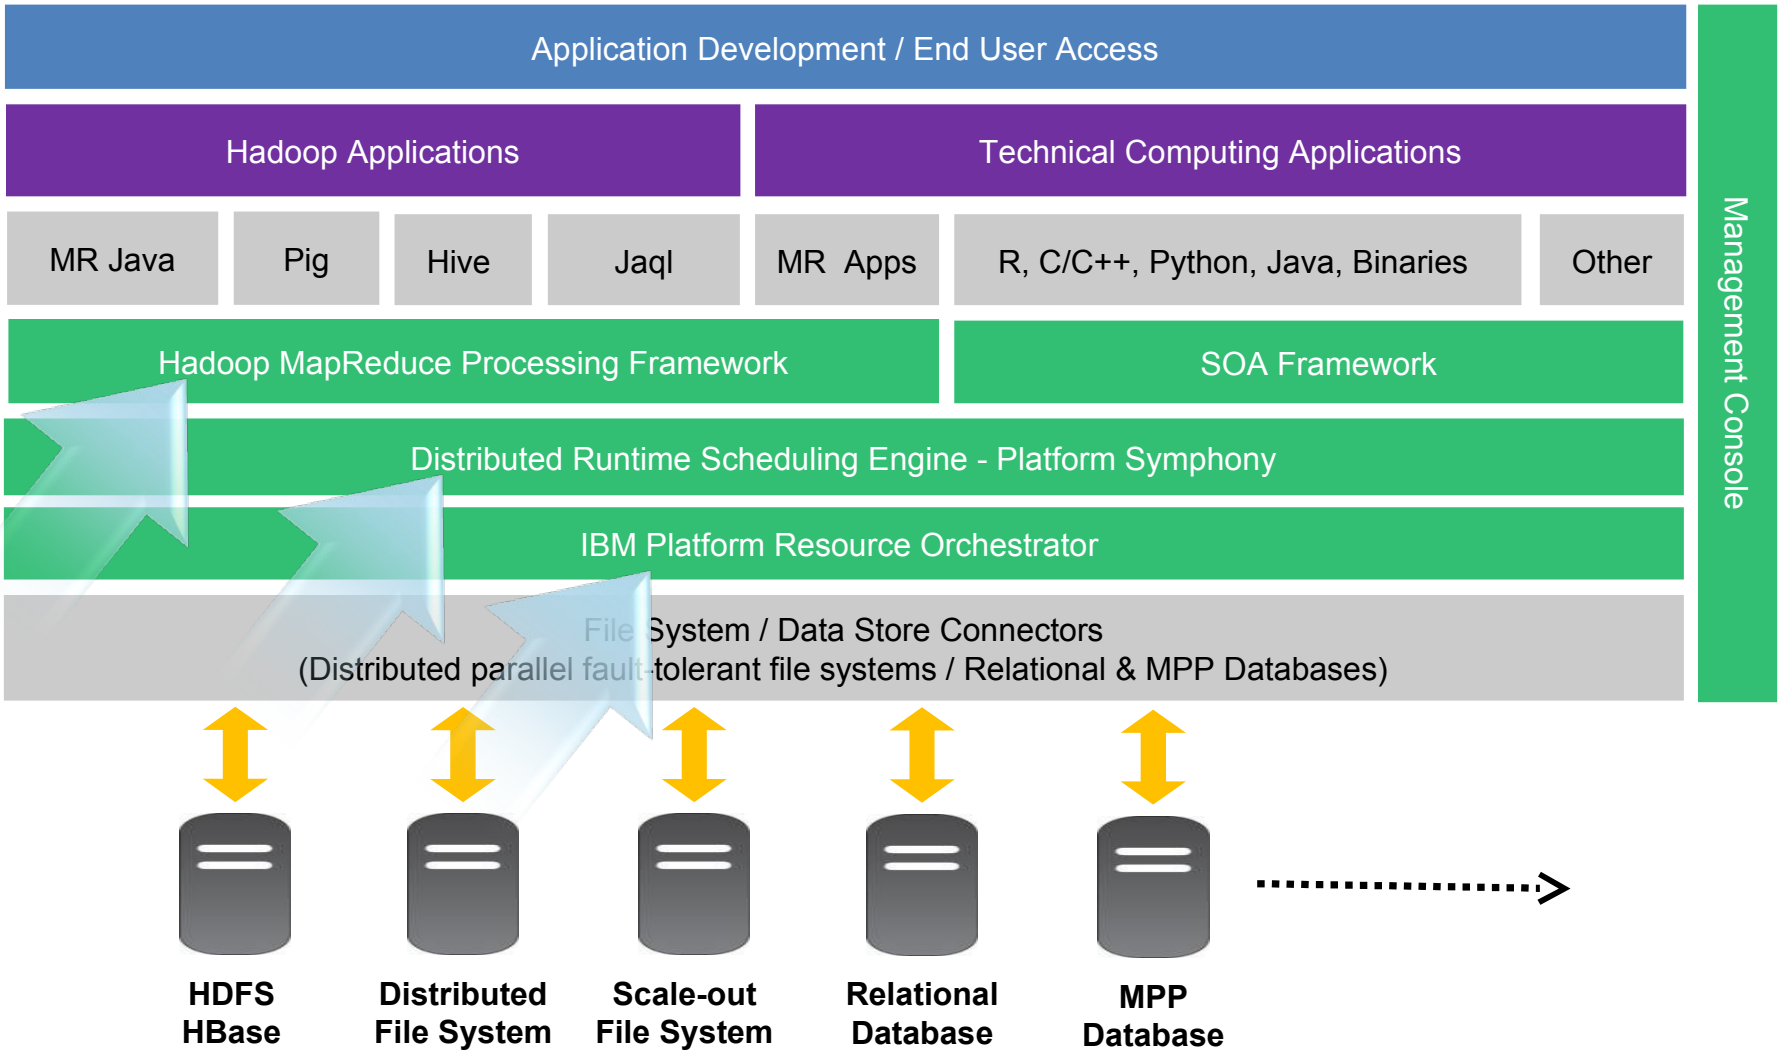

Parallel Services and MapReduce sharing an EGO Cluster

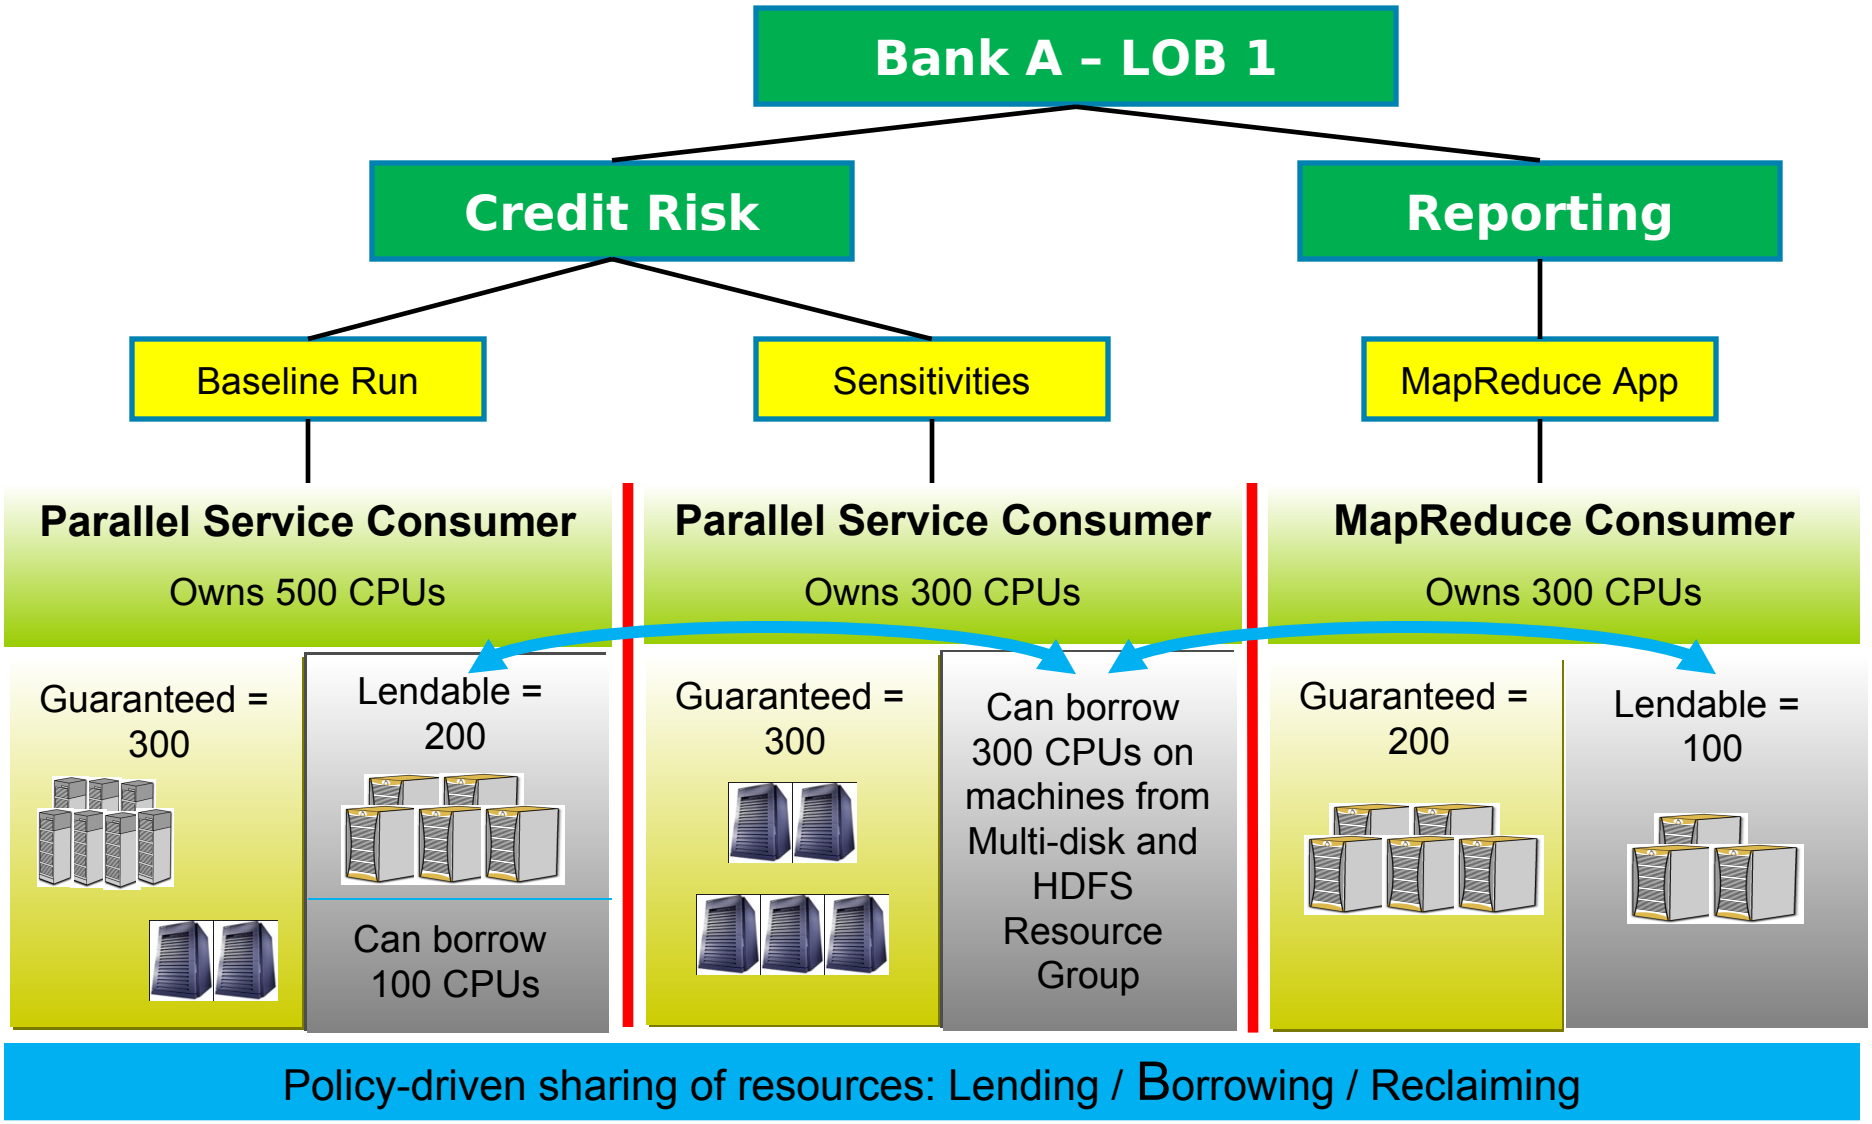

On-Demand & Reservation Allocations

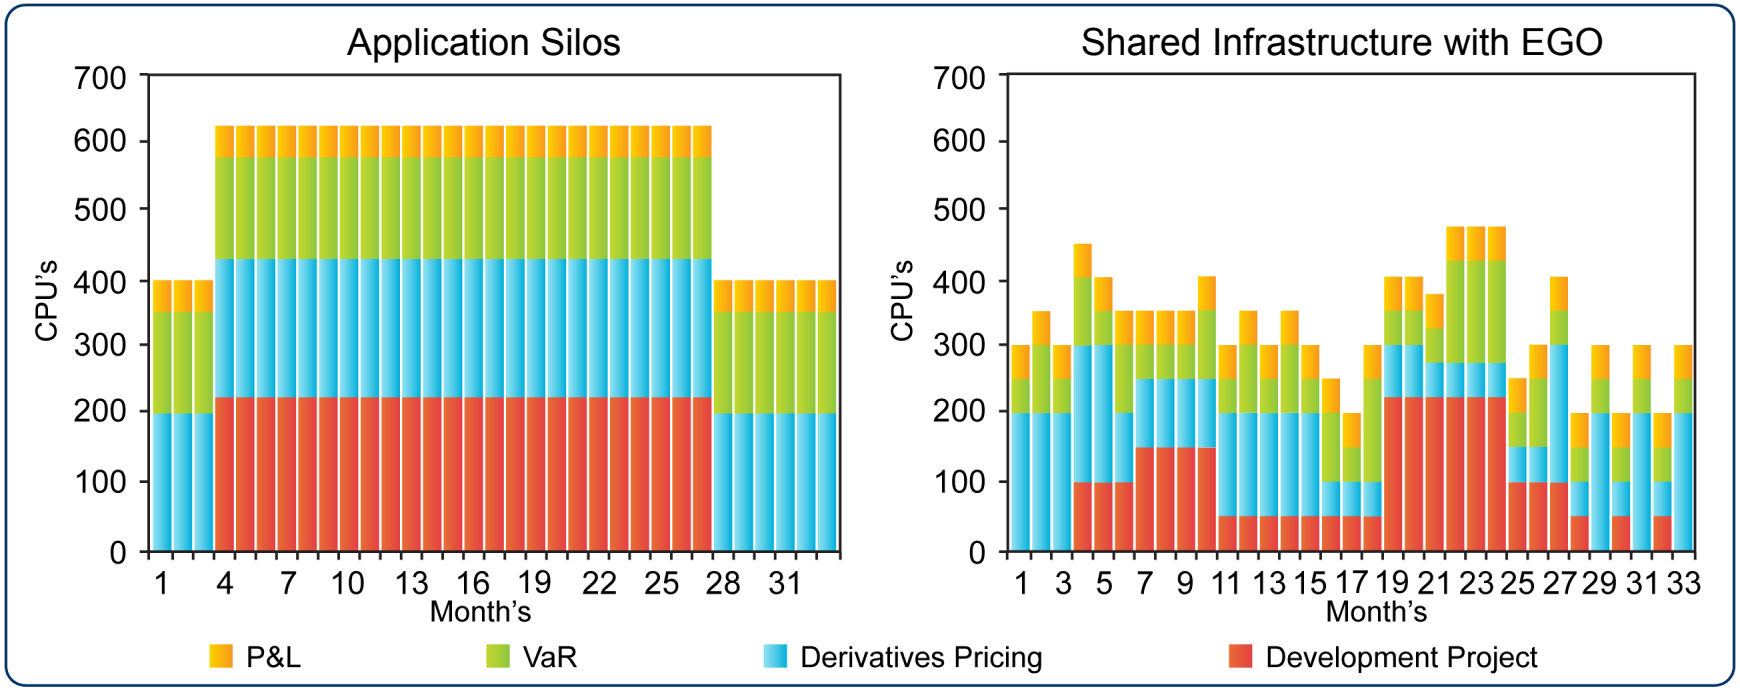

Screen Shot

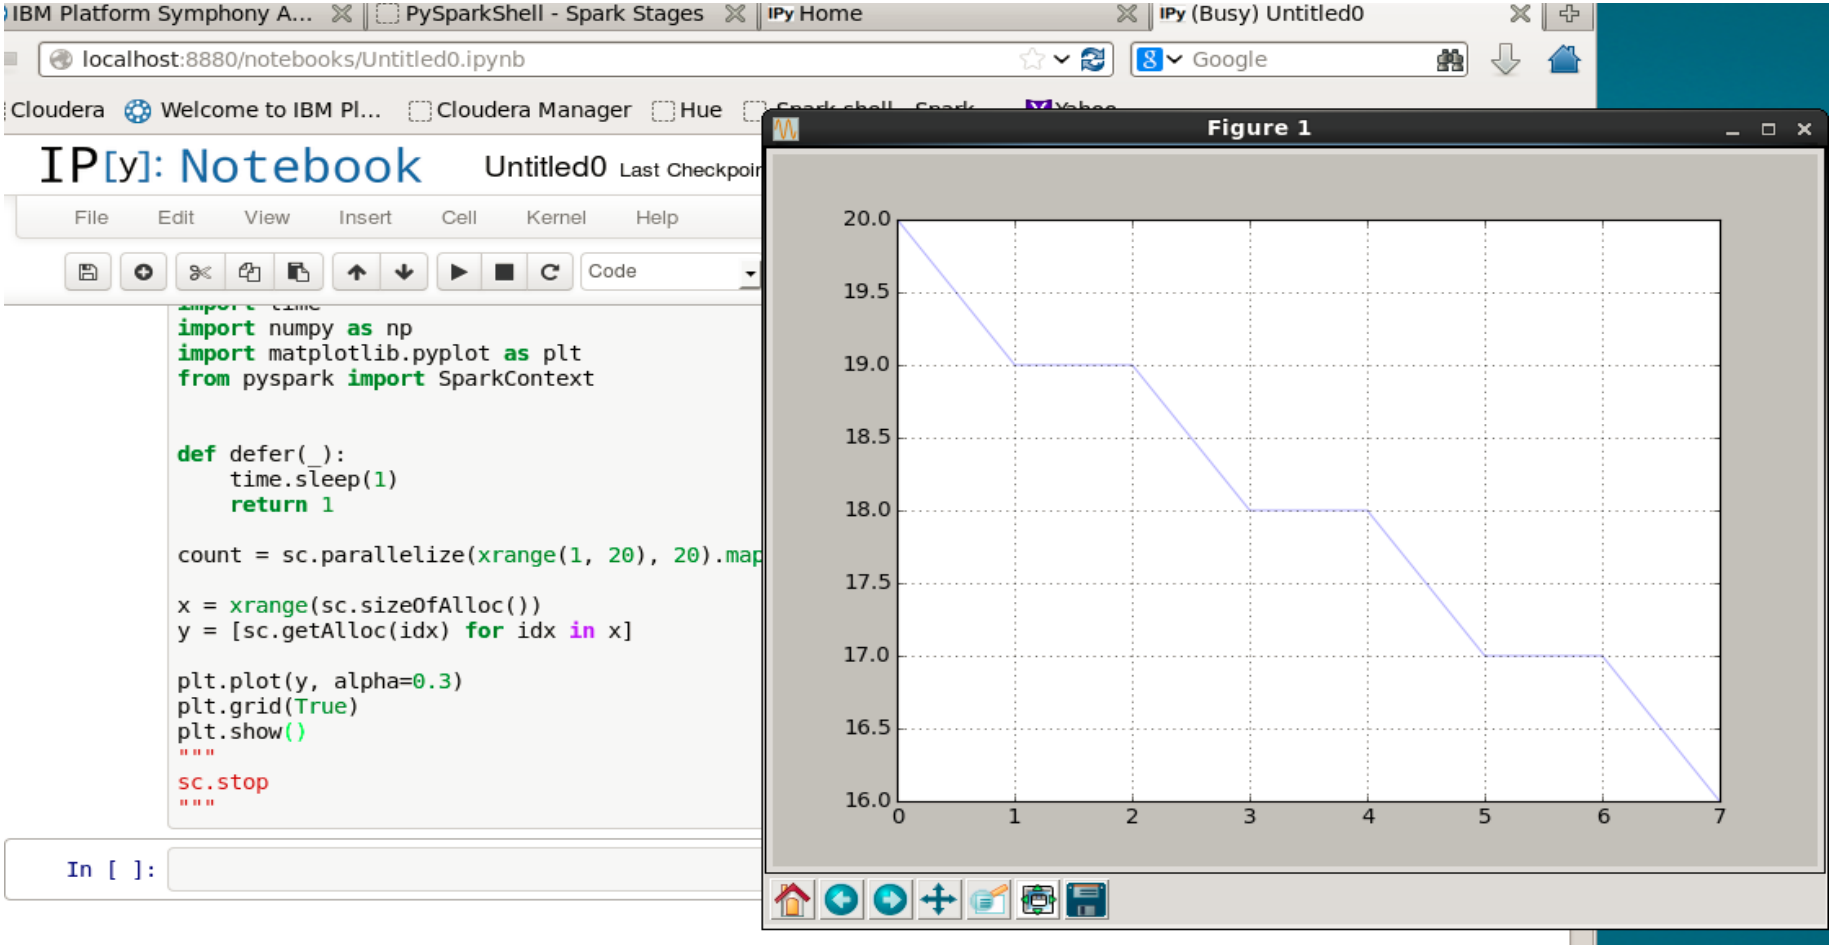

Agenda

- **Background**
- **Spark on YARN**
- **Spark on Mesos**
- **Spark on IBM Platform Symphony**
- **Spark SQL or Stream**

Spark SQL

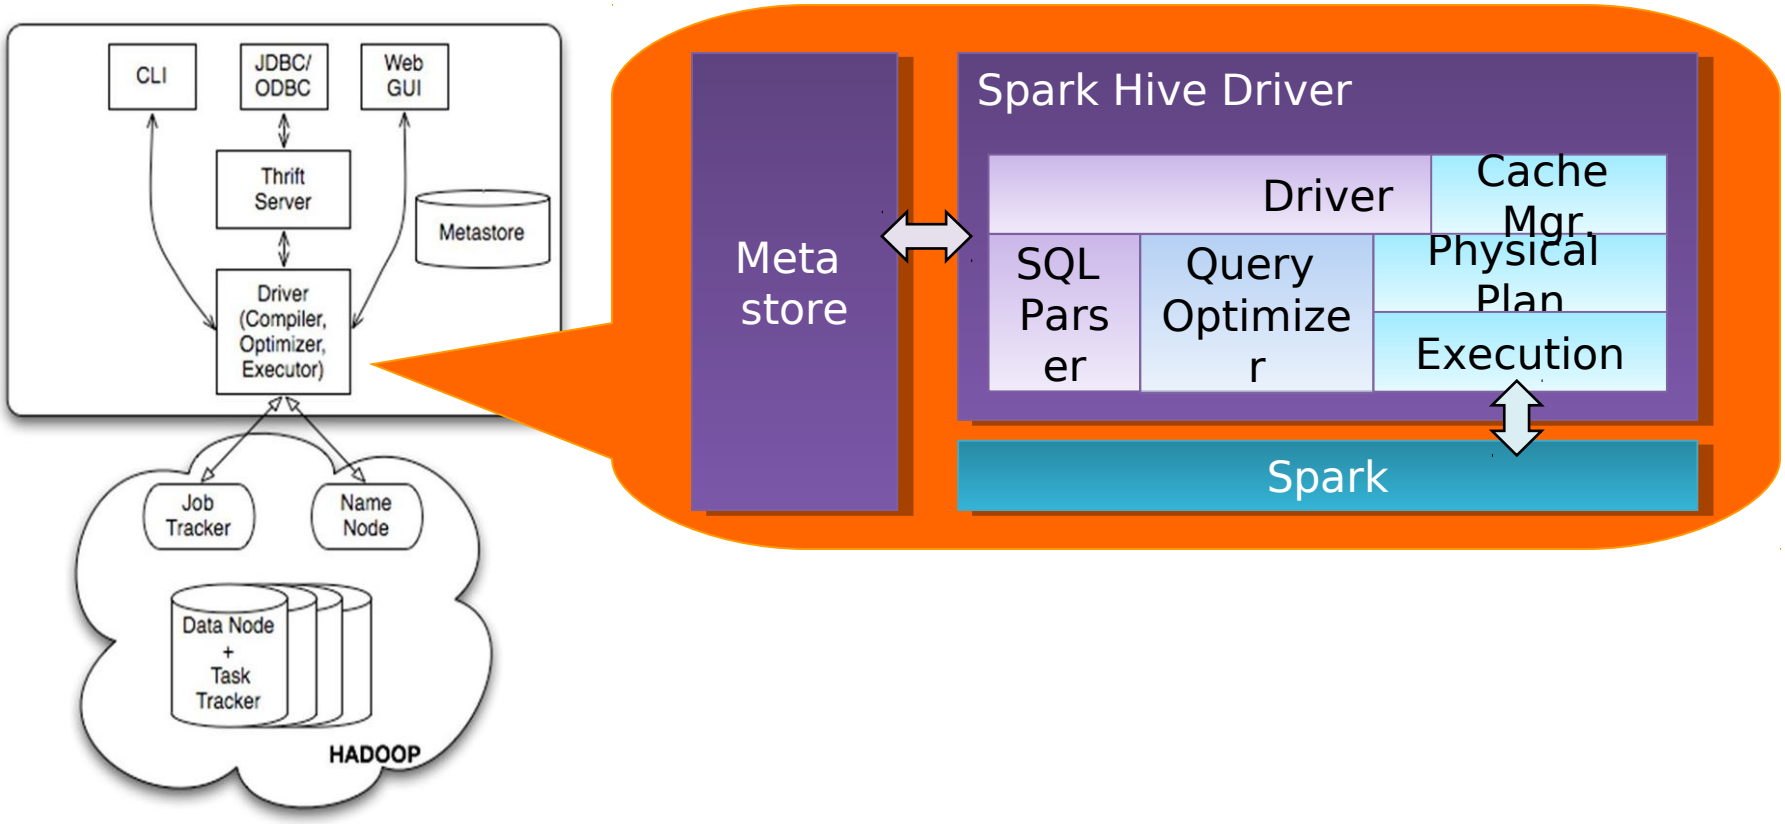

Benefits for SQL style workload

- **Centralized Scheduling**
  - Shared spark context among SQL query
  - Consolidated requests from clients
- **Resource negotiation**
  - Spark Context Scheduler
  - Request resource from RM as batch style
  - Reduce overhead to RM
- **Resource Share**
  - Better resource utilization based on workload driven
  - Reclaim happen based on priority
- **Task priority**
  - Based on existing Spark Context style