



# DATABRICKS

## **Spark SQL 1.2 Improvements and New Features**

Cheng Lian — Spark Meetup Beijing, Dec 2014

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > Enhanced Hive support
- > Misc.
- > Next steps

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > Enhanced Hive support
- > Misc.
- > Next steps

# External data source API

- > New @DeveloperApi introduced in 1.2
- > Define new input sources for Spark SQL
  - > JSON, Avro, CSV, ...
  - > Parquet, ORC, ...
  - > JDBC, C\*, HBase, ...


# External data source API

- > Mixing multiple types of data sources easily:
  - > `CREATE TEMPORARY TABLE jtable  
 USING org.apache.spark.sql.json  
 (path = "...");`
  - > `CREATE TEMPORARY TABLE ptable  
 USING org.apache.spark.sql.parquet  
 (path = "...");`
  - > `SELECT jtable.key, ptable.value  
 FROM jtable JOIN ptable  
 ON jtable.key = ptable.key;`

# External data source API

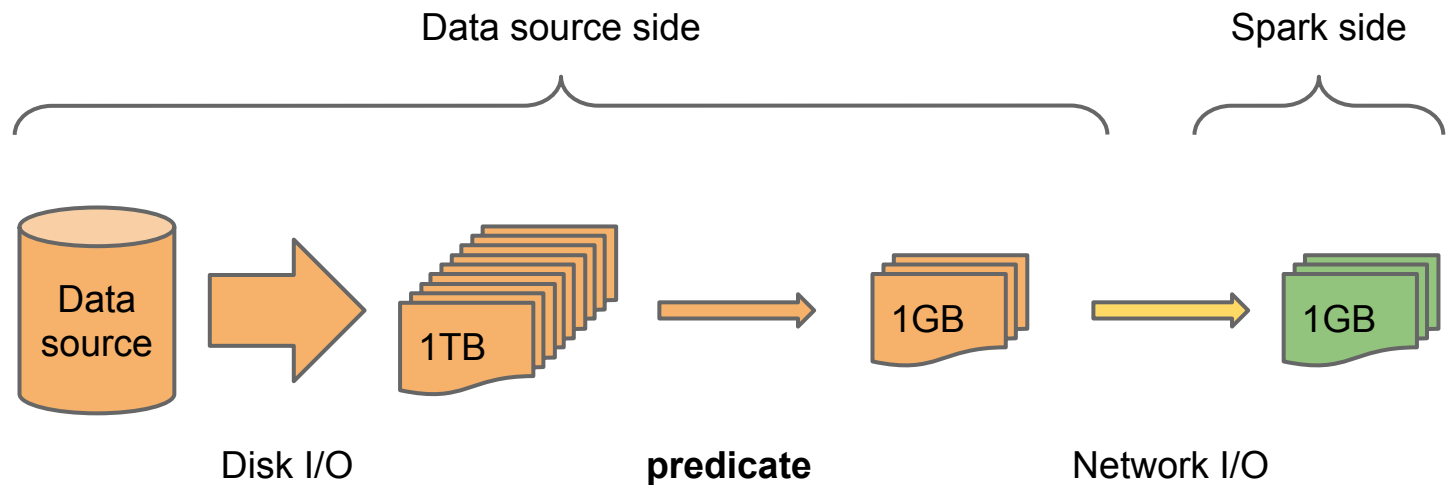
- > Supports data source specific optimizations
  - > Column pruning
  - > Pushing predicates to datasources (filter pushdown)
  - > Partition pruning (coming soon)

# External data source API



**Before...**


# External data source API



**After...**



# External data source API



**With smart data formats like Parquet and ORC,  
we can even achieve this...**

# External data source API

SELECT **Name** WHERE **ID** < 3

ID	Name	Age
1	Alice	21
2	Bob	30
3	Cart	28

Column pruning

ID	Name	Age
1	Alice	21
2	Bob	30
3	Cart	28

Predicate pushdown

# External data source API

- > Existing data sources
  - > Simple formats
    - > JSON, Avro, CSV
  - > Smart formats with column pruning and filter pushdown
    - > Parquet
    - > ORC (PR #2576 by @scwf)

# External data source API

- > Roadmap
  - > First class partitioning support with partition pruning
  - > Data sink (insertion) API
  - > Making Hive as an external data source

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > Enhanced Hive support
- > Misc.
- > Next steps

# In-memory columnar storage

- > Unified caching semantics
  - > `SchemaRDD.cache()` now uses in-memory columnar storage (finally!!!)
  - > `SQLContext.cacheTable("tbl")` is now eager by default
  - > `CACHE [LAZY] TABLE tbl [AS SELECT ...]`
- > Don't need to trigger cache materialization manually anymore
  - > `CACHE TABLE src;`  
~~`SELECT COUNT(*) FROM src;`~~

# In-memory columnar storage

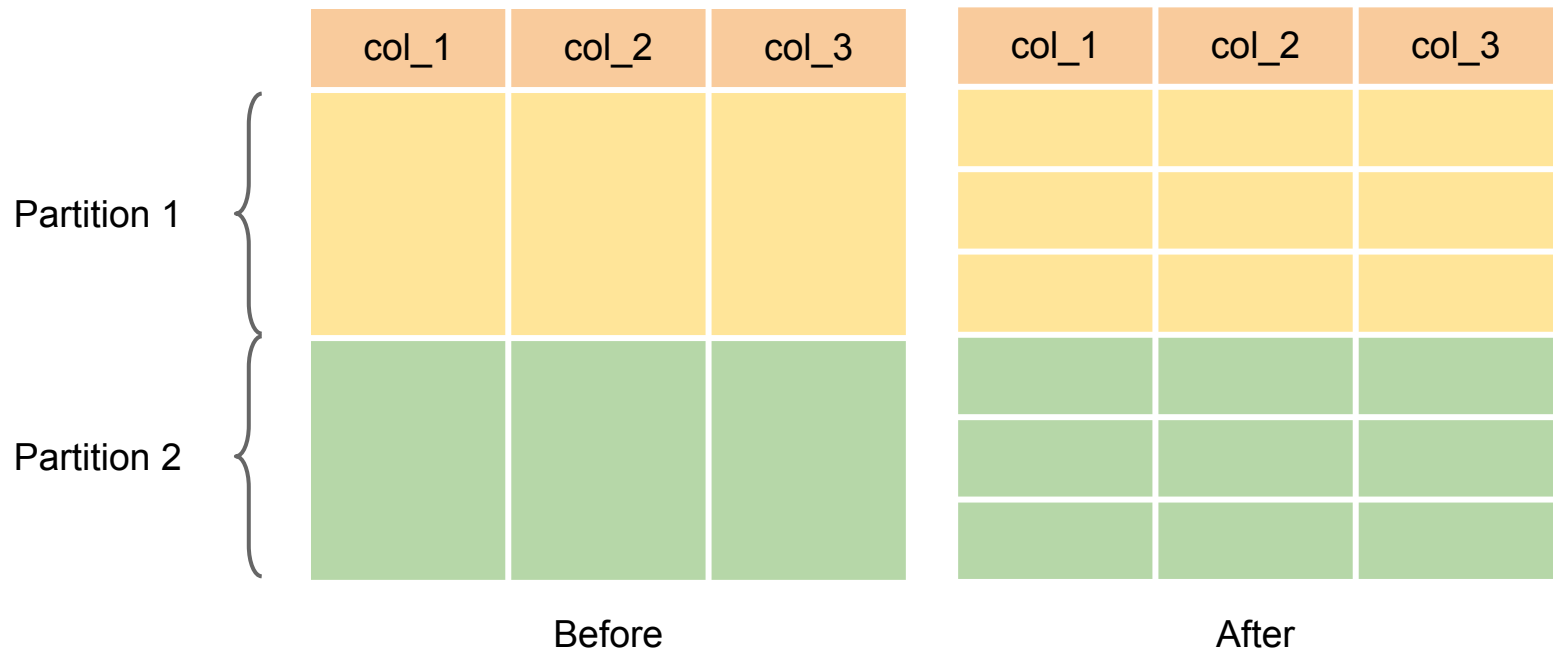
- > Query plan based cache sharing
  - > Cached queries with exactly the same result share the same underlying cache
    - > `CACHE TABLE t1 AS SELECT * FROM src;`
    - > `CACHE TABLE src;`

# In-memory columnar storage

- > Cleaner memory footprint
  - > Eliminated most boxing costs when building and accessing column buffers
  - > Introduced batched column buffer builder to avoid OOM when caching large tables



# In-memory columnar storage



# In-memory columnar storage

- > Table statistics
  - > Predicate pushdown
    - > Faster table scan
  - > Auto broadcast join
    - > Faster table join

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > **Enhanced Parquet support**
- > Enhanced Hive support
- > Misc.
- > Next steps

# Enhanced Parquet support

- > Parquet filter2 API based filter pushdown
  - > Simple predicates can be pushed down
    - > Comparisons: attr cmp literal, literal cmp attr
    - > Logical: AND, OR, NOT
  - > Enables filtering entire row groups or pages of records without touching them
  - > Unfortunately disabled by default because of PARQUET-136, expected to be re-enabled in 1.2.1

# Enhanced Parquet support

- > Ported to the data source API
  - > With simple partitioning and partition skipping support
  - > Aims to replace the old Parquet implementation

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > **Enhanced Hive support**
- > Misc.
- > Next steps

# Enhanced Hive support

- > Hive 0.13.1 support
- > Dynamic partitioning support
  - > Thanks to @baishuo from AsiaInfo!
- > View support

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > Enhanced Hive support
- > **Misc.**
- > Next steps



# Miscellaneous

- > Fixed-precision decimal type support
- > Date type support
- > UDT support (mostly for Spark ML)
- > UDF DSL
  - > 

```
val triple = (n: Int) => n * 3
```

```
table("src")  
  .select(triple.call('key) as 'k, 'value)  
  .where('key > 10)
```

# Agenda

- > External data source API
- > Enhanced in-memory columnar storage
- > Enhanced Parquet support
- > Enhanced Hive support
- > Misc.
- > **Next steps**

# Next steps

- > External data source API
  - > Data sink API
  - > First class partitioning support
- > Better way to support multiple Hive versions
- > Window function and analytics features
- > ...



DATABRICKS

**THANK YOU!**