

# The Future of Real-Time in Spark

Reynold Xin @rxin

Spark Summit, New York, Feb 18, 2016



# Why Real-Time?

Making decisions faster is valuable.

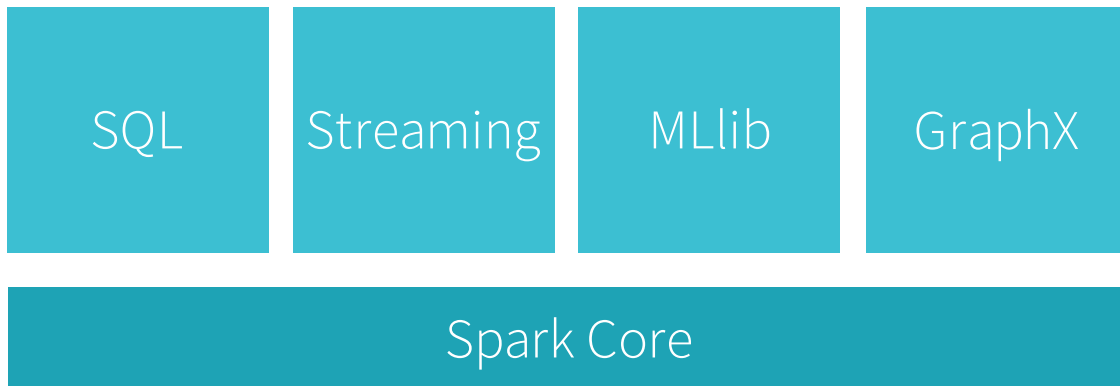
- Preventing credit card fraud
- Monitoring industrial machinery
- Human-facing dashboards
- ...

# Streaming Engine

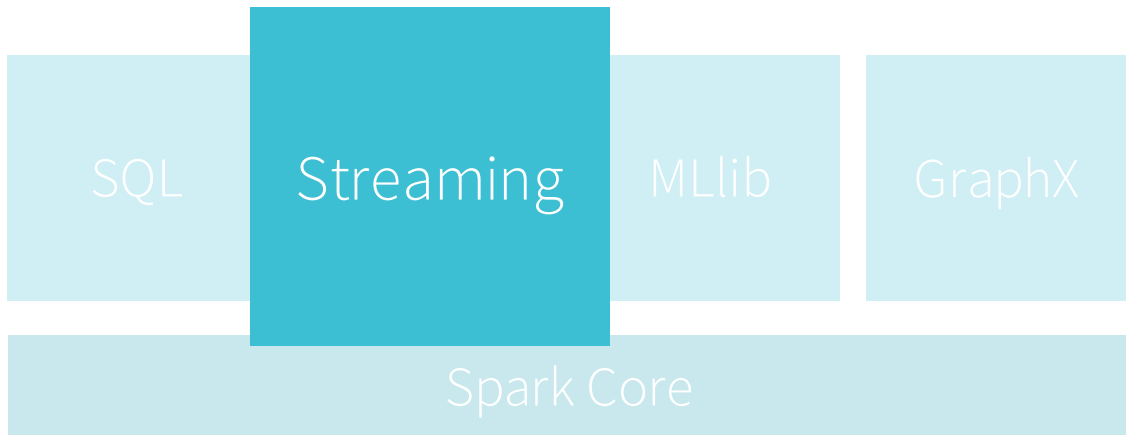
Noun.

Takes an input stream and produces an output stream.

# Spark Unified Stack



# Spark Unified Stack



Introduced 3 years ago in Spark 0.7  
50% users consider most important part of Spark


# Spark Streaming



- First attempt at unifying streaming and batch
- State management built in
- Exactly once semantics
- Features required for large clusters
  - Straggler mitigation, dynamic load balancing, fast fault-recovery

Streaming computations don't run in isolation.

# Use Case: Fraud Detection





# Continuous Application

noun.

An end-to-end application that acts on real-time data.

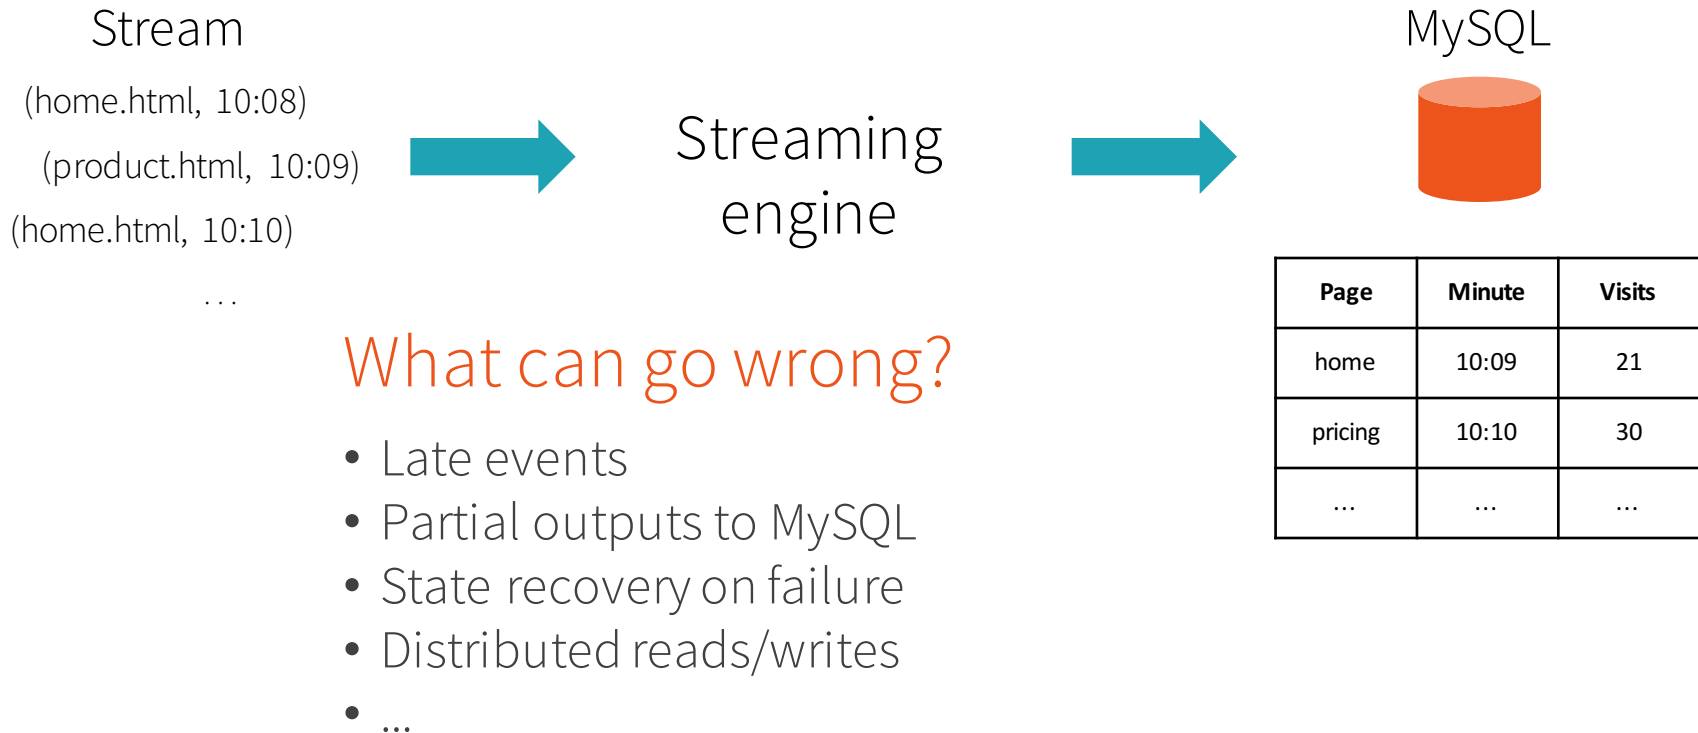
# Challenges Building Continuous Applications

Integration with non-streaming systems often an after-thought

- Interactive, batch, relational databases, machine learning, ...

Streaming programming models are complex

# Integration Example



# Complex Programming Models

## **Data**

Late arrival, varying distribution over time, ...

## **Processing**

Business logic change & new ops  
(windows, sessions)

## **Output**

How do we define  
output over time & correctness?

# Structured Streaming

The simplest way to perform streaming analytics is not having to **reason** about streaming.

[illegible][illegible]

# Single API !

# Structured Streaming

## **High-level streaming API built on Spark SQL engine**

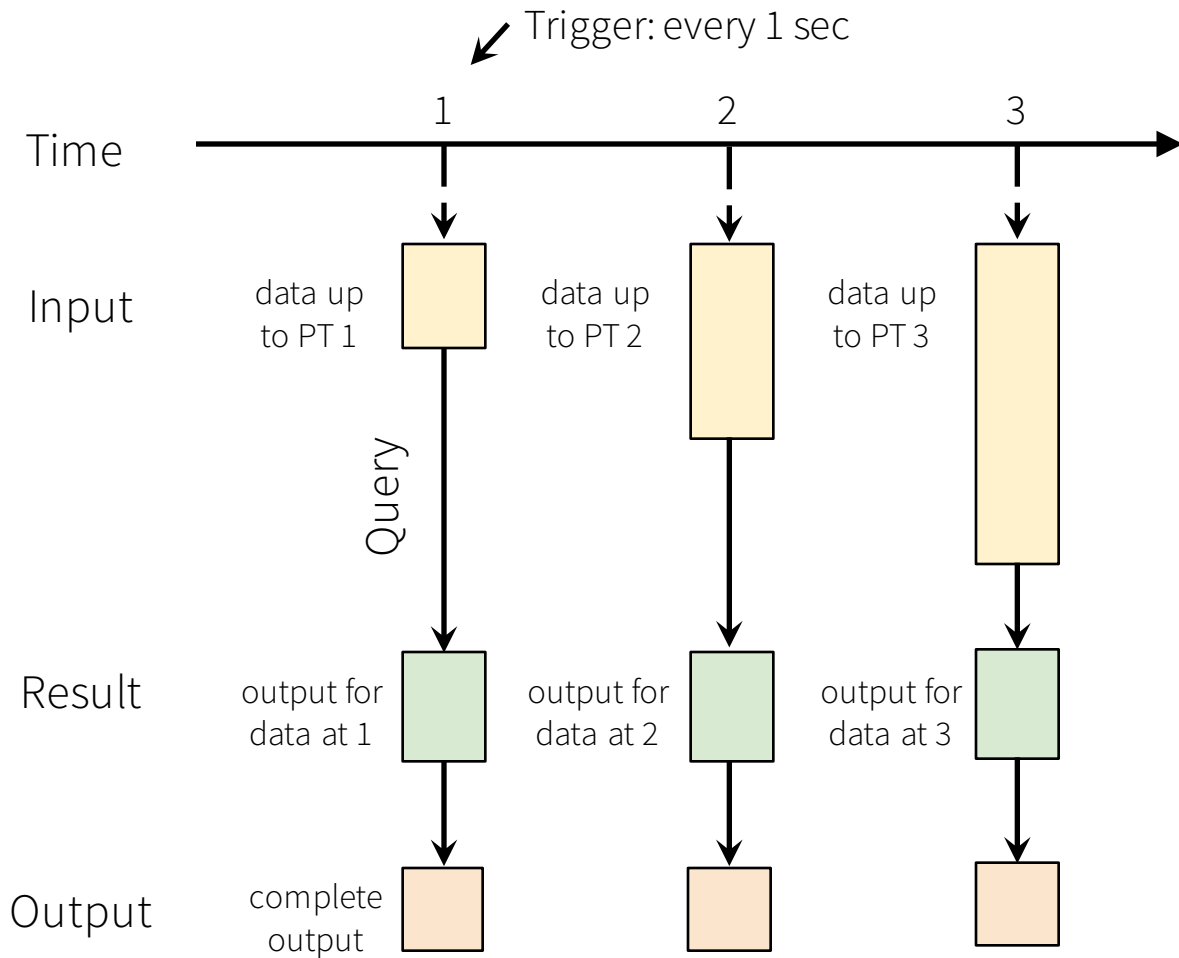
- Runs the same queries on DataFrames
- Event time, windowing, sessions, sources & sinks

## **Unifies streaming, interactive and batch queries**

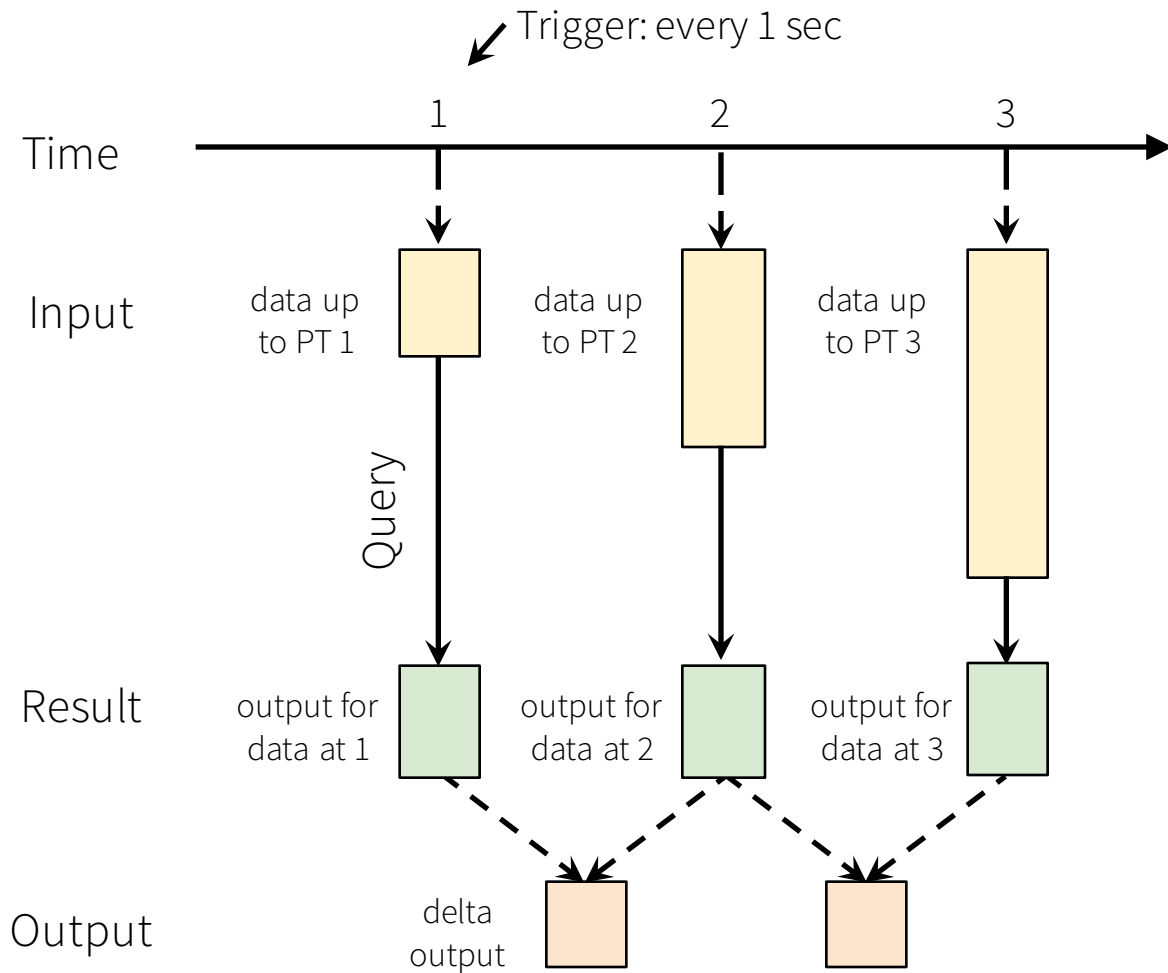
- Aggregate data in a stream, then serve using JDBC
- Change queries at runtime
- Build and apply ML models



# Model



# Model



# Model Details

**Input sources:** append-only tables

**Queries:** new operators for windowing, sessions, etc

**Triggers:** based on time (e.g. every 1 sec)

**Output modes:** complete, deltas, update-in-place

# Example: ETL

**Input:** files in S3

**Query:** map (transform each record)

**Trigger:** “every 5 sec”

**Output mode:** “new records”, into S3 sink

# Example: Page View Count

**Input:** records in Kafka

**Query:** `select count(*) group by page, minute(evtime)`

**Trigger:** “every 5 sec”

**Output mode:** “update-in-place”, into MySQL sink

**Note:** this will automatically update “old” records on late data!

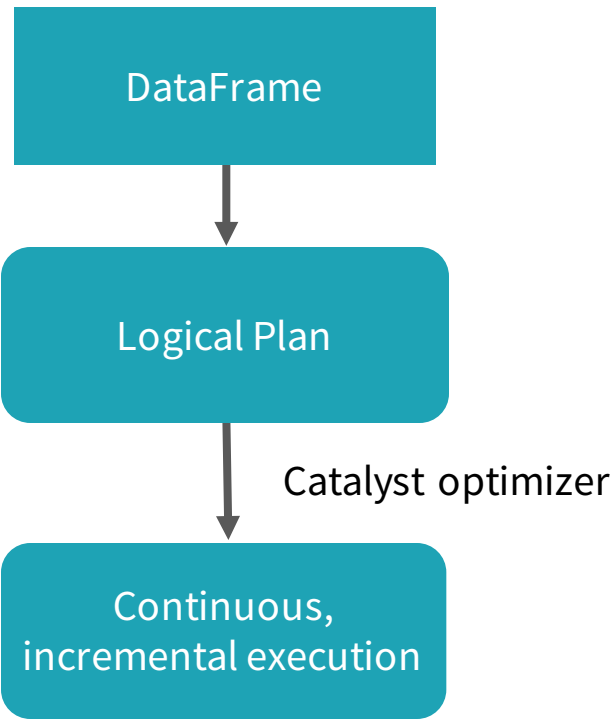
# Execution

## **Logically:**

DataFrame operations on static data  
(i.e. as easy to understand as batch)

## **Physically:**

Spark automatically runs the query in  
streaming fashion  
(i.e. incrementally and continuously)



# Example: Batch Aggregation

```
logs = ctx.read.format("json").open("s3://logs")
```

```
logs.groupBy(logs.user_id).agg(sum(logs.time))  
  .write.format("jdbc")  
  .save("jdbc:mysql://...")
```


# Example: Continuous Aggregation

```
logs = ctx.read.format("json").stream("s3://logs")
```

```
logs.groupBy(logs.user_id).agg(sum(logs.time))  
    .write.format("jdbc")  
    .stream("jdbc:mysql://...")
```



# Automatic Incremental Execution



# Rest of Spark will follow

- Interactive queries should just work
- Spark's data source API will be updated to support seamless streaming integration
  - Exactly once semantics **end-to-end**
  - Different output modes (complete, delta, update-in-place)
- ML algorithms will be updated too


# What can we do with this that's hard with other engines?

Ad-hoc, interactive queries

Dynamic changing queries

Benefits of Spark: elastic scaling, straggler mitigation, etc

# Use Case: Fraud Detection



# Timeline

## Spark 2.0

- API foundation
- Kafka, file systems, and databases
- Event-time aggregations

## Spark 2.1 +

- Continuous SQL
- BI app integration
- Other streaming sources / sinks
- Machine learning

Thank you.

@rxin

